

CUSTOMER SERVICE AGENTS

Hope or enforcement: who keeps each promise?

One fictional shop's support agent, built three ways and simulated. Every design has a rule for everything.

A rule the model keeps is hope

A limit the agent cannot reach is a boundary.
Only the second is a control.

How much of each policy is hope

Design 1, one agent, two connectors: 97%

Design 2, a harness of tools: 73%

Design 3, gated narrow agents: 23%

What one run can reach


Design 1

- 38,000 customer records
- Unlimited refunds
- Any email address

Design 3

- One verified customer
- 100 GBP per order
- No other address

Design 3 turns one prompt into a flow



The specialists' tools take no customer argument.

As reach falls, the policy gets smaller

It no longer has to say "do not delete the database" to an agent that has no way to.

The policy shows which boundary to build next

New address: confirm by email link first

Refunds stop at 250 GBP per customer in 30 days

Claims need an image

Two vouchers per customer in 90 days

6 of 8

client promises kept by a boundary in Design 3

Design 1 keeps none by a boundary, Design 2 one. Accurate, polite replies rest on a model in every design.

Show the client the record, not a score

A promise is only as strong as the weakest thing it rests on.

sgit.ai/articles/hope-or-enforcement.html